

# experiments1

prateek sharma

05 September 2011

## Contents

<b>1</b>	<b>Page Sharing</b>	<b>2</b>
1.1	Scanning vs Disk based sharing . . . . .	2
<b>2</b>	<b>KSM</b>	<b>2</b>
2.1	Exp 1: KSM effectiveness . . . . .	2
<b>3</b>	<b>Analysis of shared pages</b>	<b>2</b>
3.1	Exp 2: Pages shared by flag type . . . . .	2
3.2	Exp 2.1 : KSM with no pagecache pages . . . . .	2
3.3	Exp 2.2: Ftrace Overhead for KSM . . . . .	2
3.4	Exp 3: Page sharing over time . . . . .	3
<b>4</b>	<b>Lookahead optimization</b>	<b>3</b>
4.1	Exp 4: Lookahead success . . . . .	3
4.2	Exp 5: Substrings in shared-map. . . . .	3
<b>5</b>	<b>Problem of double-caching</b>	<b>3</b>
5.1	Exp 6: Memory savings with exclusive caches . . . . .	3
5.2	Exp 7: Overhead of ksm-exclusive-cache . . . . .	3
<b>6</b>	<b>Qualitative survey of dynamic memory management for VMs.</b>	<b>4</b>
6.1	Exp 8: Memory mountains . . . . .	4

## 1 Page Sharing

Survey of everything. Some history?

### 1.1 Scanning vs Disk based sharing

2 competing approaches. (2)

## 2 KSM

Diagram of operations, history, and some implementation details Focus on generic nature

### 2.1 Exp 1: KSM effectiveness

Small experiment which shows that KSM shares large % of same pages. Done earlier in fingerprinting project. Expected result : > 90% sharing of KSM, so good enough. Reason : establish some ground truths: KSM works . Setup : Random workload (doesn't matter) and take fingerprint and see KSM shar% . Run with 1,2,3 VMs. (1)

## 3 Analysis of shared pages

### 3.1 Exp 2: Pages shared by flag type

Run some benchmarks (static VMs just booted up ; Kernbench ; HTTP-perf) and see what kinds of pages are shared by flag type. Reason: establish some ground truths : sharing is feasible . Also answer : what kinds of pages are shared? Setup: (1,2,3,5) VMs with different OS running same benchmarks. (diff VMs ; same kernel ; same /var/www)

### 3.2 Exp 2.1 : KSM with no pagecache pages

Run KSM but skip all guest pagecache pages.

### 3.3 Exp 2.2: Ftrace Overhead for KSM

Since we are recording all KSM events (just for few experiments only) whats the overhead of that? 1GB/minute data collected.

### 3.4 Exp 3: Page sharing over time

Run some benchmarks and record pages shared over the duration of benchmark. Also record **KSM overhead**. Reason: Show that KSM overhead is significant enough, thus implying the need for some optimizations. Setup: (2) VMs running benchmarks. KSM being profiled using perf.

(1)

## 4 Lookahead optimization

### 4.1 Exp 4: Lookahead success

Run benchmarks on VMs (1,2,3) to on and record lookahead successes. Also record **KSM overhead** **Compare vanilla KSM overhead with lookahead-optimization**

### 4.2 Exp 5: Substrings in shared-map.

Record consecutive pages being shared in some benchmarks. Reason : justify why lookahead works. Setup: tracedump analysis simple python script

(3)

## 5 Problem of double-caching

### 5.1 Exp 6: Memory savings with exclusive caches

How many pages are there in both places? Setup : Benchmarks on VMs.

### 5.2 Exp 7: Overhead of ksm-exclusive-cache

Run benchmarks on VMs to record KSM overhead (with ex cache) Reason : scanning vast host page cache could be significant overhead. Also savings might help.

Some caching theory references. (1)

## 6 Qualitative survey of dynamic memory management for VMs.

### 6.1 Exp 8: Memory mountains

Look at this problem like L1/2 cache , and build mem-mountains in these cases: (normal ; no guest cache ; no host cache ; swap as ramdisk ) Setup : could use IOZone or randal bryant's simple program. Reason : Demonstrate the latencies/throughput of various caches. **This depends on lots of factors like IO schedulers, FS, virtual disk layout etc. Do for any one, for now.**